

Introduction to Machine Learning

New York University
Summer 2020

Homework 1

Venkata Naga Sai Kiran Challa (Net ID: vc2118)

1. Let $\{x_1, x_2, \dots, x_n\}$ be a set of points in d -dimensional space. Suppose we wish to produce a single point estimate $\mu \in R^d$ that minimizes the *mean squared-error*:

$$\frac{1}{n}(\|x_1 - \mu\|_2^2 + \|x_2 - \mu\|_2^2 + \dots + \|x_n - \mu\|_2^2)$$

Find a closed form expression for μ and prove that your answer is correct.

Answer:

First, lets try out with all the possible values that we can come up with on an example and then try to prove it.

Let's consider an example vector like, $[2, 2, 3, 4, 5]$. So, $x_1 = 2$, $x_2 = 2$, $x_3 = 3$ and so on.

Right out of the bat we can think that 2 which occurred the most can minimise it. So, let see.

$$meansquared - error = \frac{1}{5}((2-2)^2 + (2-2)^2 + (3-2)^2 + (4-2)^2 + (5-2)^2) = 2.8$$

Now, it seems pretty close to minimum value 2 , but we will try rest of the values. Let's try 5, the maximum value.

$$meansquared - error = \frac{1}{5}((2-5)^2 + (2-5)^2 + (3-5)^2 + (4-5)^2 + (5-5)^2) = 4.6$$

It's a large value compare to the previous. As we are at it let's check the rest of the values as well. We checked the minimum value (which is also the *mode* of the vector) and maximum values. Let's check with the mean and median of the set as well.

We can easily find out the mean and median of the set as,

$$\text{Mean} = \frac{1}{5}(2 + 2 + 3 + 4 + 5) = 3.2$$

Median = Middle value = 3

Therefore,

$$\text{meansquared - error} = \frac{1}{5}((2-3.2)^2 + (2-3.2)^2 + (3-3.2)^2 + (4-3.2)^2 + (5-3.2)^2) = 1.36$$

$$\text{meansquared - error} = \frac{1}{5}((2-3)^2 + (2-3)^2 + (3-3)^2 + (4-3)^2 + (5-3)^2) = 1.4$$

Out of all the possible values we see that *mean* value minimises the most. Is it just for this case or is a general property? Let's find out.

We know that when we take differentiation *once* and equate to zero, we get the point that maximises/minimises the expression and if the *second* differentiation is greater than zero, that point gives us the minimum and vice versa.

Let \bar{X} be the vector of values $\{x_1, x_2, \dots, x_n\}$.

$$\Rightarrow \text{L2 Norm of } \bar{X} = \|X\|_2 = \sqrt{\sum_{i=1}^n x_i^2} = x \cdot x^T$$

Thus, we can convert the *mean squared-error* (**MSE**) as,

$$\begin{aligned} \text{MSE} &= \frac{1}{n}[(\bar{X} - \mu) \cdot (\bar{X} - \mu)^T] \\ &= \frac{1}{n} \sum_{i=1}^n [(x_i - \mu) \cdot (x_i - \mu)^T] \\ &= \frac{1}{n} \sum_{i=1}^n [x_i \cdot x_i^T - x_i \mu^T - \mu x_i^T + \mu^2] \end{aligned} \tag{1}$$

We know from the properties of vector spaces that, $x_i \mu^T = \mu x_i^T$

$$\begin{aligned} \text{MSE} &= \frac{1}{n} \sum_{i=1}^n [x_i \cdot x_i^T - 2\mu x_i^T + \mu^2] \\ &= \frac{1}{n} \left[\sum_{i=1}^n x_i \cdot x_i^T - 2\mu \sum_{i=1}^n x_i + \sum_{i=1}^n \mu^2 \right] \\ &= \frac{1}{n} \left[\sum_{i=1}^n x_i \cdot x_i^T - 2\mu \sum_{i=1}^n x_i + n\mu^2 \right] \end{aligned} \tag{2}$$

This simplified equation can be *differentiated w.r.t μ* ,

$$\Rightarrow \frac{1}{n}[-2 \sum_{i=1}^n x_i + 2.n\mu] = 0 \quad (3)$$

$$\Rightarrow \boxed{\mu = \frac{\sum_{i=1}^n x_i}{n}} \quad (4)$$

We can note that the *second* differentiation of the above equation (3) is,

$$\begin{aligned} \text{MSE}' &= -2 \sum_{i=1}^n x_i + 2.n\mu \\ \text{MSE}'' &= 2.n > 0 \end{aligned} \quad (5)$$

Hence, we can justify that above equation (4) will minimises the *mean squared-error*.

2. Not all norms behave the same; for instance, the ℓ_1 -norm of a vector can be dramatically different from the ℓ_2 -norm, especially in high dimensions. Prove the following norm inequalities for d-dimensional vectors, starting from the definitions provided in class and lecture notes. (Use any algebraic technique/result you like, as long as you cite it.)
 - a. $\|x\|_\infty \leq \|x\|_2 \leq \sqrt{d}\|x\|_\infty$
 - b. $\|x\|_\infty \leq \|x\|_1 \leq d\|x\|_\infty$

Answer:

I will be taking generalised approach to this problem. First I am going to prove the first half in both *a* and *b*, and latter on proceed with proving the second half. We know,

$$\begin{aligned} \ell_1 - \text{norm} &= \|X\|_1 = \sum_{i=1}^d |x_i| \\ \ell_2 - \text{norm} &= \|X\|_2 = \sqrt{\sum_{i=1}^d x_i^2} \\ \ell_p - \text{norm} &= \|X\|_p = \left(\sum_{i=1}^d x_i^p \right)^{\frac{1}{p}} \\ \ell_\infty - \text{norm} &= \|X\|_\infty = \lim_{x \rightarrow \infty} \left(\sum_{i=1}^d x_i^p \right)^{\frac{1}{p}} \\ &= \max_i |x_i|, \forall i = 1, 2, ..d \end{aligned} \quad (6)$$

Let's look into *infinity norm*, ℓ_∞ – norm. Why is it, $\max_i |x_i|$? Well if you think about it, as $p \rightarrow \infty$, we can realise that most dominant term in x_i over takes the rest of the values. For more formal definition and proof, you can follow the Wikipedia's page, by clicking [here](#).

For comparing different norm values, let's see them on a unit circle, taking $x_1 = [-1, 1]$ and $x_2 = [-1, 1]$ and restricting our space between $[0, 1], [1, 0], [0, -1]$ and $[-1, 0]$. Thus, we will be finding x_1 and x_2 such that,

$$\begin{aligned}\ell_1 - \text{norm} &= \|X\|_1 = 1 \\ \ell_2 - \text{norm} &= \|X\|_2 = 1 \\ \ell_\infty - \text{norm} &= \|X\|_\infty = 1\end{aligned}\tag{7}$$

$\ell_1 - \text{norm}$				
x_1	1	0.8	0.6	0.5
x_2	0	0.2	0.4	0.5
$\ell_2 - \text{norm}$				
x_1	1	0.8	0.4	0.5
x_2	0	0.6	0.91	0.87
$\ell_\infty - \text{norm}$				
x_1	1	0		
x_2	0	1		

Table 1: Change of x_2 w.r.t x_1 .

Thus,

$$\begin{aligned}\ell_1 = \|X\|_1 = \|(x_1, x_2)\| &\iff |x_1| + |x_2| = 1 \\ &\implies \text{In } Q1, x_2 = 1 - x_1 \\ &\implies \text{In } Q2, x_2 = 1 + x_1 \\ &\implies \text{In } Q3, x_2 = -x_1 - 1 \\ &\implies \text{In } Q4, x_2 = x_1 - 1\end{aligned}\tag{8}$$

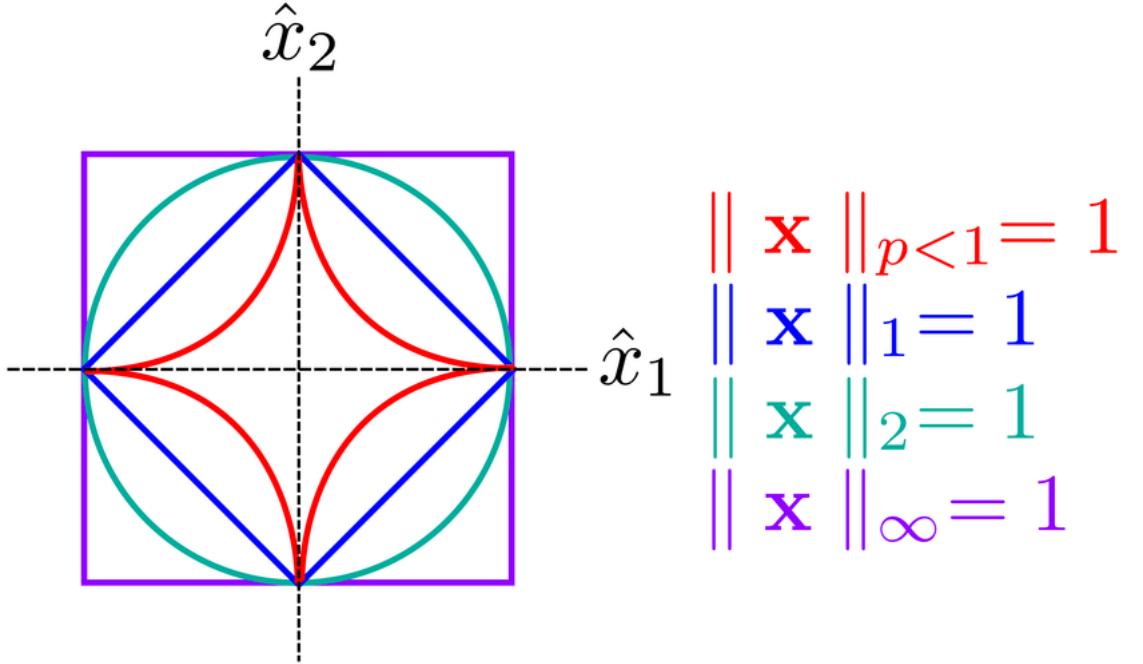
We, got a square which is rotated 45° , i.e a rhombus.(unit circle for ℓ_1 norm)
For ℓ_2 ,

$$\ell_2 = \|X\|_2 = \left(\sum_{i=1}^d x_i^2 \right)^{\frac{1}{2}} \iff x_1^2 + x_2^2 = 1\tag{9}$$

This is an equation of a circle with radius one.(unit circle for ℓ_2 norm)
For ℓ_∞ ,

$$\begin{aligned}
\ell_\infty &= \max_i\{|1|, |x_2|\} \quad \forall x_2 \leq 1 \implies x_1 = 1 \\
&= \max_i\{|-1|, |x_2|\} \quad \forall x_2 \leq 1 \implies x_1 = -1 \\
&= \max_i\{|x_1|, |1|\} \quad \forall x_1 \leq 1 \implies x_2 = 1 \\
&= \max_i\{|x_1|, |-1|\} \quad \forall x_1 \leq 1 \implies x_2 = -1
\end{aligned} \tag{10}$$

This gives us a square (unit circle for ℓ_∞ norm). They are shown in the below figure.



Source: Research gate

Figure 1 : Plot of different norms.

We can notice that unit circle for ℓ_1 norm is less than ℓ_2 norm, which indeed is less than ℓ_∞ norm. This tells us,

$$\Rightarrow \|x\|_1 \geq \|x\|_2 \geq \|x\|_\infty \tag{11}$$

Why is that so? Because the unit circle for $\|x\|_1$ norm is smaller than the unit circle for $\|x\|_2$ norm, i.e, for same values of x_1 and x_2 , we get larger value for ℓ_1 compared to ℓ_2 . Similarly, this applies for ℓ_∞ norm as well. For example,

For $x_1 = 0.6$ and $x_2 = 0.3$,

$$\begin{aligned}
\|x\|_1 &= 0.6 + 0.3 = 0.9, \text{ and} \\
\|x\|_2 &= \sqrt{0.6^2 + 0.3^2} = 0.67 \\
\|x\|_\infty &= \max_i\{0.6, 0.3\} = 0.6
\end{aligned} \tag{12}$$

This, helped us to complete the first half of the question, which told us to prove $\|x\|_\infty \leq \|x\|_2$ and $\|x\|_\infty \leq \|x\|_1$.

For the second half, instead of me proving for d and \sqrt{d} individually, we will take a generalised approach.

We know, that below equation is true as we approximated the ℓ_∞ norm.

$$|x_i| \leq \|x\|_\infty, \forall i = \{1, 2, \dots, d\}$$

Now let's raise the power of the above equation to p . As, $1 \leq p < \infty$ the sign of inequality does not change.

$$|x_i|^p \leq \|x\|_\infty^p, \forall i = \{1, 2, \dots, d\}$$

We have this equation for every $i = \{1, 2, \dots, d\}$, i.e

$$\begin{aligned} |x_1|^p &\leq \|x\|_\infty^p \\ |x_2|^p &\leq \|x\|_\infty^p \\ &\vdots \\ |x_d|^p &\leq \|x\|_\infty^p \end{aligned} \tag{13}$$

Adding all the terms in the above equation set (13), we get the below equation. Notice here that d gets multiplied in the RHS as there are d possible values of x_i .

$$|x_1|^p + |x_2|^p + \dots + |x_d|^p \leq d \cdot \|x\|_\infty^p$$

Taking the p^{th} root and simplifying further.

$$\begin{aligned} (|x_1|^p + |x_2|^p + \dots + |x_d|^p)^{\frac{1}{p}} &\leq (d \cdot \|x\|_\infty^p)^{\frac{1}{p}} \\ \left(\sum_{i=1}^d |x_i|^p\right)^{\frac{1}{p}} &\leq (d \cdot \|x\|_\infty^p)^{\frac{1}{p}} \\ \implies \|x\|_p &\leq d^{\frac{1}{p}} (\|x\|_\infty^p)^{\frac{1}{p}} \\ \implies \|x\|_p &\leq d^{\frac{1}{p}} \cdot \|x\|_\infty \end{aligned} \tag{14}$$

Now, if we substitute $p = 1$ for $\|x\|_1$ norm and $p = 2$ for $\|x\|_2$ norm. We get,

$$\boxed{\|x\|_1 \leq d \cdot \|x\|_\infty \text{ and } \|x\|_2 \leq \sqrt{d} \cdot \|x\|_\infty}$$

Combining equation (11) and **above**, we obtain our answers.

$$\boxed{\|x\|_\infty \leq \|x\|_2 \leq \sqrt{d} \|x\|_\infty, \text{ and, } \|x\|_\infty \leq \|x\|_1 \leq d \|x\|_\infty}$$

Assignment: Exploratory data analysis

TODO: Edit this cell to fill in your NYU Net ID and your name:

- **Net ID:** vc2118
- **Name:** Venkata Naga Sai Kiran Challa

Introduction

In this assignment, we will practice using exploratory data analysis on Google's COVID-19 Community Mobility data.

This data was collected from Google Maps users around the world over the last few months - including you, *if* you have Google Maps on your phone and have turned on the Location History setting. It combines location history from a large number of users to capture the overall increase or decrease in time spent in places such as: retail and recreation facilities, groceries and pharmacies, parks, transit stations, workplaces, and residences.

As you work through this notebook, you will see that some text and code cells are marked with a "TODO" at the top. You'll have to edit these cells to fill in the code or answer the questions as indicated.

Learn about the data

First, it is worthwhile to learn more about the data: how it is collected, what is included, how Google gets consent to collect this data, and how user privacy is protected. Google provides several resources for learning about the data:

- [Blog post](#)
- [About this data](#)
- [Understand the data](#)

Read in data

Now you are ready to read the data into your notebook.

Visit Google's web page for the [COVID-19 Community Mobility](#) project to get the URL for the data.

(Specific instructions will depend on your browser and operating system, but on my laptop, I can get the URL by right-clicking on the button that says "Download global CSV" and choosing "Copy Link Address".)

Then, in the following cells, use that URL to read the data into a pandas Data Frame called `df`. (You can follow the example in the "Exploratory data analysis" notebook from this week's lesson.)

In [1]:

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

/usr/local/lib/python3.6/dist-packages/statsmodels/tools/_testing.py:19: FutureWarning: pandas.util.testing is deprecated. Use the functions in the public API at pandas.testing instead.
    import pandas.util.testing as tm
```

Note Pdf is not so aesthetically pleasing to look at. So I recommend to checkout my github link for the same notebook [here](#)

In [2]:

```
# TODO Q1
url = 'https://www.gstatic.com/covid19/mobility/Global_Mobility_Report.csv?cachebust=694ae9957380f150'
df = pd.read_csv(url)

/usr/local/lib/python3.6/dist-packages/IPython/core/interactiveshell.py:2718: DtypeWarning:
  Columns (3) have mixed types. Specify dtype option on import or set low_memory=False.
    interactivity=interactivity, compiler=compiler, result=result)
```

Use the `info()` and `head()` functions to show some basic information about the data and to look at the first few samples.

In [3]:

```
# TODO Q2
# use info()
df.info() # Gives us basic information of the dataset
# It's 516697 * 13, and has null values
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 516697 entries, 0 to 516696
Data columns (total 13 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   country_region_code  515810 non-null   object 
 1   country_region      516697 non-null   object 
 2   sub_region_1        501307 non-null   object 
 3   sub_region_2        298720 non-null   object 
 4   iso_3166_2_code     187299 non-null   object 
 5   census_fips_code    298834 non-null   float64
 6   date               516697 non-null   object 
 7   retail_and_recreation_percent_change_from_baseline 400610 non-null   float64
 8   grocery_and_pharmacy_percent_change_from_baseline 385086 non-null   float64
 9   parks_percent_change_from_baseline       238643 non-null   float64
 10  transit_stations_percent_change_from_baseline 275427 non-null   float64
 11  workplaces_percent_change_from_baseline 503228 non-null   float64
 12  residential_percent_change_from_baseline 284348 non-null   float64
dtypes: float64(7), object(6)
memory usage: 51.2+ MB
```

In [4]:

```
# TODO Q3
# use head()
df.head() # Prints first 5 rows
```

Out [4]:

	country_region_code	country_region	sub_region_1	sub_region_2	iso_3166_2_code	census_fips_code	date	retail_and_rec
0	AE	United Arab Emirates	NaN	NaN	NaN	NaN	2020-02-15	
1	AE	United Arab Emirates	NaN	NaN	NaN	NaN	2020-02-16	
2	AE	United Arab Emirates	NaN	NaN	NaN	NaN	2020-02-17	
3	AE	United Arab Emirates	NaN	NaN	NaN	NaN	2020-02-18	
4	AE	United Arab Emirates	NaN	NaN	NaN	NaN	2020-02-19	

Basic data manipulations

The data includes a date field, but it may have been read in as a string, rather than as a `datetime`. If that's the case, use `to_datetime()` to convert the field into a datetime format. (You can follow the example in the

“Exploratory data analysis” notebook from this week’s lesson.)

Then, use `info()` again to make sure your change was applied. Note the difference in the output, relative to the cell above.

In [5]:

```
# TODO Q4
# df['date'] = ...
# Converting to date format
df['date'] = pd.to_datetime(df['date'])
df.info()    # Check if its converted - yes, it is converted
```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 516697 entries, 0 to 516696
Data columns (total 13 columns):
 # Column Non-Null Count Dtype
--- --
 0 country_region_code 515810 non-null object
 1 country_region 516697 non-null object
 2 sub_region_1 501307 non-null object
 3 sub_region_2 298720 non-null object
 4 iso_3166_2_code 187299 non-null object
 5 census_fips_code 298834 non-null float64
 6 date 516697 non-null datetime64[ns]
 7 retail_and_recreation_percent_change_from_baseline 400610 non-null float64
 8 grocery_and_pharmacy_percent_change_from_baseline 385086 non-null float64
 9 parks_percent_change_from_baseline 238643 non-null float64
 10 transit_stations_percent_change_from_baseline 275427 non-null float64
 11 workplaces_percent_change_from_baseline 503228 non-null float64
 12 residential_percent_change_from_baseline 284348 non-null float64
dtypes: datetime64[ns] (1), float64(7), object(5)
memory usage: 51.2+ MB

Next, you are going to extract the subset of data for the location of your choice. You can choose any location **except Brooklyn, New York**. (You can’t choose Brooklyn because the example code I’m about to show you is for Brooklyn.)

The data is reported for different regions, with different levels of granularity available. This is best explained by example:

Suppose I want the overall trend from the entire U.S. I would use the subset of data where `country_region` is equal to “United States” and `sub_region_1` is null:

```
df_subset = df[(df['country_region'].eq("United States")) & (df['sub_region_1'].isnull())]
```

Suppose I want the overall trend from the entire state of New York: I would use the subset of data where `country_region` is equal to “United States”, `sub_region_1` is equal to “New York”, and `sub_region_2` is null:

```
df_subset = df[(df['country_region'].eq("United States")) & (df['sub_region_1'].eq("New York")) & (df['sub_region_2'].isnull())]
```

Suppose I want the overall trend from Brooklyn, New York (Kings County): I would use the subset of data where `country_region` is equal to “United States”, `sub_region_1` is equal to “New York”, and `sub_region_2` is equal to “Kings County”:

```
df_subset = df[(df['country_region'].eq("United States")) & (df['sub_region_1'].eq("New York")) & (df['sub_region_2'].eq("Kings County"))]
```

In the following cell(s), fill in the code to create a data frame `df_subset` with data from a single location. You

can go down to the `sub_region_1` level or the `sub_region_2` level - depending on the location you chose, the finer level of granularity may not be available.

In [6]:

```
# TODO Q5
# df_subset =
#df[(df['country_region'].eq("India"))]
df[(df['country_region'].eq("India")) & (df['sub_region_1'].eq("Maharashtra"))].info() # No Mumbai, so we can stop here - no sub_region_2 values
df_subset = df[(df['country_region'].eq("India")) & (df['sub_region_1'].eq("Maharashtra"))]
```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 114 entries, 93711 to 93824
Data columns (total 13 columns):
 # Column Non-Null Count Dtype

 0 country_region_code 114 non-null object
 1 country_region 114 non-null object
 2 sub_region_1 114 non-null object
 3 sub_region_2 0 non-null object
 4 iso_3166_2_code 114 non-null object
 5 census_fips_code 0 non-null float64
 6 date 114 non-null datetime64[ns]
 7 retail_and_recreation_percent_change_from_baseline 114 non-null float64
 8 grocery_and_pharmacy_percent_change_from_baseline 114 non-null float64
 9 parks_percent_change_from_baseline 114 non-null float64
 10 transit_stations_percent_change_from_baseline 114 non-null float64
 11 workplaces_percent_change_from_baseline 114 non-null float64
 12 residential_percent_change_from_baseline 114 non-null float64
dtypes: datetime64[ns] (1), float64(7), object(5)
memory usage: 12.5+ KB

Is the data complete, or is some data not available for the location you have chosen? In the following cell, write code to check for missing data in the `...percent_change_from_baseline` fields.

In [7]:

```
# TODO Q6
# df_subset
# Kind of re-iterating the above step, check if there are missing values
df_subset.info() # There are no missing values except for sub_region_2 and census_fips_code, and all the percentage changes are intact
```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 114 entries, 93711 to 93824
Data columns (total 13 columns):
 # Column Non-Null Count Dtype

 0 country_region_code 114 non-null object
 1 country_region 114 non-null object
 2 sub_region_1 114 non-null object
 3 sub_region_2 0 non-null object
 4 iso_3166_2_code 114 non-null object
 5 census_fips_code 0 non-null float64
 6 date 114 non-null datetime64[ns]
 7 retail_and_recreation_percent_change_from_baseline 114 non-null float64
 8 grocery_and_pharmacy_percent_change_from_baseline 114 non-null float64
 9 parks_percent_change_from_baseline 114 non-null float64
 10 transit_stations_percent_change_from_baseline 114 non-null float64
 11 workplaces_percent_change_from_baseline 114 non-null float64
 12 residential_percent_change_from_baseline 114 non-null float64
dtypes: datetime64[ns] (1), float64(7), object(5)
memory usage: 12.5+ KB

In [8]:

```
# Having a look at data
df_subset
```

Out [8] :

country_region_code	country_region	sub_region_1	sub_region_2	iso_3166_2_code	census_fips_code	date	retail_and_recreation_percent_change_from_baseline
93711	IN	India	Maharashtra	NaN	IN-MH	NaN	2020-02-15
93712	IN	India	Maharashtra	NaN	IN-MH	NaN	2020-02-16
93713	IN	India	Maharashtra	NaN	IN-MH	NaN	2020-02-17
93714	IN	India	Maharashtra	NaN	IN-MH	NaN	2020-02-18
93715	IN	India	Maharashtra	NaN	IN-MH	NaN	2020-02-19
...
93820	IN	India	Maharashtra	NaN	IN-MH	NaN	2020-06-03
93821	IN	India	Maharashtra	NaN	IN-MH	NaN	2020-06-04
93822	IN	India	Maharashtra	NaN	IN-MH	NaN	2020-06-05
93823	IN	India	Maharashtra	NaN	IN-MH	NaN	2020-06-06
93824	IN	India	Maharashtra	NaN	IN-MH	NaN	2020-06-07

114 rows × 13 columns

TODO Q7: Edit this cell to answer the following question: Is the data complete, or is some relevant data missing? Why would some locations only have partial data available (missing some `...percent_change_from_baseline` fields for some dates)? Include a short quote from the material you read in the “Learn about the data” section to answer this question.

In [0] :

```
#Added empty code block for easy finding the answer.
```

Q7 Response: Fortunately, the subset I chose did not have any missing values except for `sub_region_2` and `census_fips_code`. Although, having missing values are perfectly acceptable as data that doesn't meet Google's privacy threshold is not included in the dataset, and this data should be considered as unknown data rather than assuming that those values are not busy. Source : https://support.google.com/covid19-mobility/answer/9825414?hl=en&ref_topic=9822927

In [0] :

```
#End of the answer!!
```

For this data, the `date` field is important, but we don't necessarily care about the absolute date. Instead, we care about how many days have elapsed since the first confirmed case of COVID-19 in this location, how many days have elapsed since a “stay at home” order or similar rule was established in this location (if there was one) and how many days have elapsed since it was lifted (if applicable).

For example, in Brooklyn, New York, I might compute:

```
days_since_lockdown = (df_ny_all['date'] - pd.to_datetime('2020-03-20 00:00:00')).dt.days.values  
# NYC lockdown March 20, 2020 https://www.nytimes.com/2020/03/20/us/coronavirus-today.html
```

Compute “days since [some relevant COVID-19 date]” for your location. In a comment, explain the significance of the date you have chosen, and include a link to a news article or other reference supporting the significance of the date. (The news article does not have to be in English.)

In [0]:

```
# TODO Q8
days_since_lockdown = (df_subset[df_subset['date'] >= pd.to_datetime('2020-03-25 00:00:00')]['date'] - pd.to_datetime('2020-03-25 00:00:00')).dt.days # days from the lockdown initiation
# India lockdown 25th March, 2020 https://www.nytimes.com/2020/03/24/world/asia/india-coronavirus-lockdown.html and is still in effect
# Eventhough India has extended the lockdown 3 times, I chose that date as it's the first lockdown and people will be more willing to follow rules.
```

In [10]:

```
days_since_lockdown # Index and value, I have created the data set down before plotting
```

Out[10]:

```
93750      0
93751      1
93752      2
93753      3
93754      4
..
93820     70
93821     71
93822     72
93823     73
93824     74
Name: date, Length: 75, dtype: int64
```

In [11]:

```
days_since_lockdown = days_since_lockdown.values # We need this for plotting
days_since_lockdown
```

Out[11]:

```
array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14, 15, 16,
       17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31, 32, 33,
       34, 35, 36, 37, 38, 39, 40, 41, 42, 43, 44, 45, 46, 47, 48, 49, 50,
       51, 52, 53, 54, 55, 56, 57, 58, 59, 60, 61, 62, 63, 64, 65, 66, 67,
       68, 69, 70, 71, 72, 73, 74])
```

Visualize data

Finally, we are going to visualize the changes in human mobility over this time.

In the following cell, create a figure with six subplots, arranged vertically. (You can refer to the example in the “Python + numpy” notebook from this week’s lesson.) On the horizontal axis, put the `days_since...` array you computed in the previous cell. On the vertical axes, show:

- `retail_and_recreation_percent_change_from_baseline` in the top subplot
- `grocery_and_pharmacy_percent_change_from_baseline` in the next subplot
- `parks_percent_change_from_baseline` in the next subplot
- `transit_stations_percent_change_from_baseline` in the next subplot
- `workplaces_percent_change_from_baseline` in the next subplot
- `residential_percent_change_from_baseline` in the bottom subplot

In [12]:

```
# TODO Q9
#First create the subset of data after lockdown
```

```
df_after = df_subset[df_subset['date'] >= pd.to_datetime('2020-03-25 00:00:00')]  
df_after # We got data from 25th March
```

Out [12] :

country_region_code	country_region	sub_region_1	sub_region_2	iso_3166_2_code	census_fips_code	date	retail_and_recreation_percent_change_from_baseline
93750	IN	India	Maharashtra	NaN	IN-MH	NaN	2020-03-25
93751	IN	India	Maharashtra	NaN	IN-MH	NaN	2020-03-26
93752	IN	India	Maharashtra	NaN	IN-MH	NaN	2020-03-27
93753	IN	India	Maharashtra	NaN	IN-MH	NaN	2020-03-28
93754	IN	India	Maharashtra	NaN	IN-MH	NaN	2020-03-29
...
93820	IN	India	Maharashtra	NaN	IN-MH	NaN	2020-06-03
93821	IN	India	Maharashtra	NaN	IN-MH	NaN	2020-06-04
93822	IN	India	Maharashtra	NaN	IN-MH	NaN	2020-06-05
93823	IN	India	Maharashtra	NaN	IN-MH	NaN	2020-06-06
93824	IN	India	Maharashtra	NaN	IN-MH	NaN	2020-06-07

75 rows x 13 columns

In [0] :

```
#Let's Prepare  
x = (df_subset['date'] - pd.to_datetime('2020-03-25 00:00:00')).dt.days.values # We will  
plot from begining to get an idea.  
y_retail_and_recreation_percent_change_from_baseline = df_subset['retail_and_recreation_p  
ercent_change_from_baseline']  
y_grocery_and_pharmacy_percent_change_from_baseline = df_subset['grocery_and_pharmacy_per  
cent_change_from_baseline']  
y_parks_percent_change_from_baseline = df_subset['parks_percent_change_from_baseline']  
y_transit_stations_percent_change_from_baseline = df_subset['transit_stations_percent_cha  
nge_from_baseline']  
y_workplaces_percent_change_from_baseline = df_subset['workplaces_percent_change_from_b  
aseline']  
y_residential_percent_change_from_baseline = df_subset['residential_percent_change_from_b  
aseline']
```

In [34] :

```
#Plot plot plot  
  
plt.style.use("seaborn") # To add looks  
plt.subplots_adjust(left=0.125, bottom=0.1, right=0.9, top=2.9, wspace=0.2, hspace=0.4) #  
To provide spacing between plots  
# Set up a subplot grid that has height 6 and width 1,  
# and set the first such subplot as active.  
plt.subplot(6, 1, 1)  
  
# Make the first plot  
plt.plot(x, y_retail_and_recreation_percent_change_from_baseline)  
plt.title('Retail and Recreation')  
plt.xlabel('Days since lockdown')  
plt.ylabel('% change')
```

```

# Set the second subplot as active, and make the second plot.
plt.subplot(6, 1, 2)
plt.plot(x, y_grocery_and_pharmacy_percent_change_from_baseline)
plt.title('Grocery and Pharmacy')
plt.xlabel('Days since lockdown')
plt.ylabel('% change')

# Set the third subplot as active, and make the third plot.
plt.subplot(6, 1, 3)
plt.plot(x, y_parks_percent_change_from_baseline)
plt.title('Parks')
plt.xlabel('Days since lockdown')
plt.ylabel('% change')

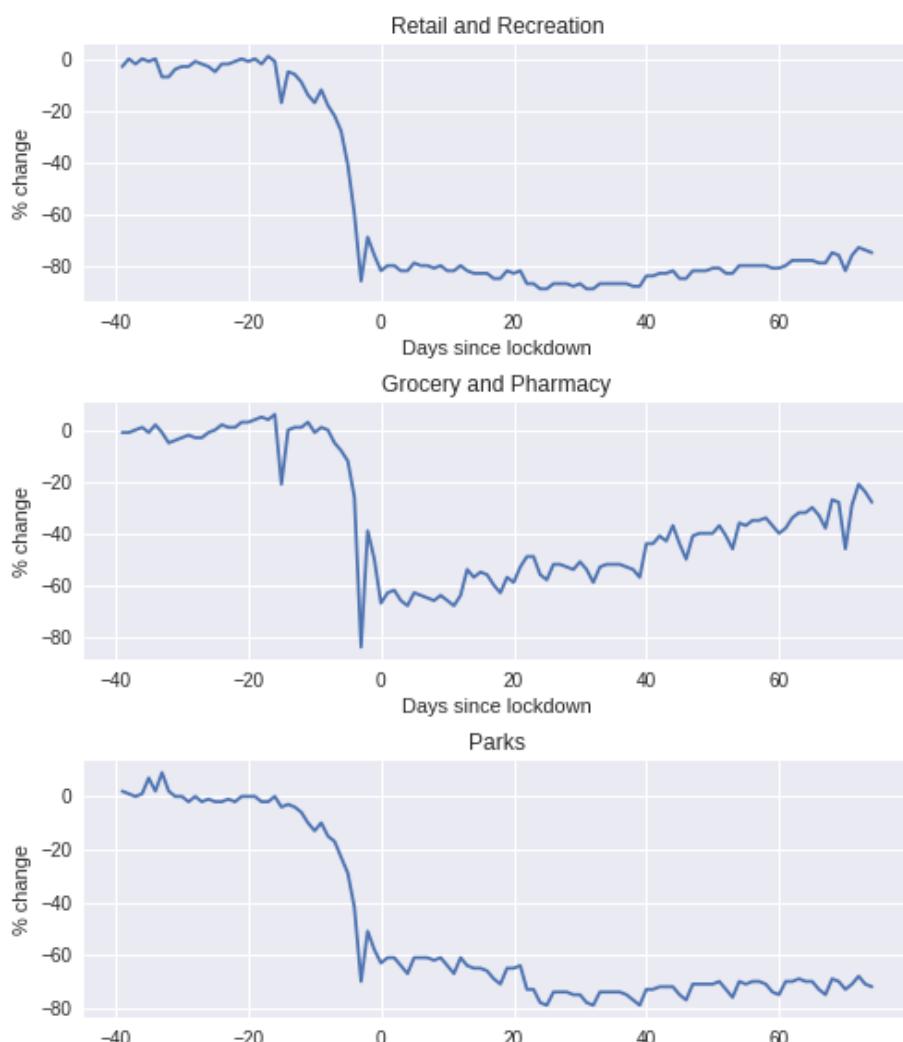
# Set the fourth subplot as active, and make the fourth plot.
plt.subplot(6, 1, 4)
plt.plot(x, y_transit_stations_percent_change_from_baseline)
plt.title('Transit stations')
plt.xlabel('Days since lockdown')
plt.ylabel('% change')

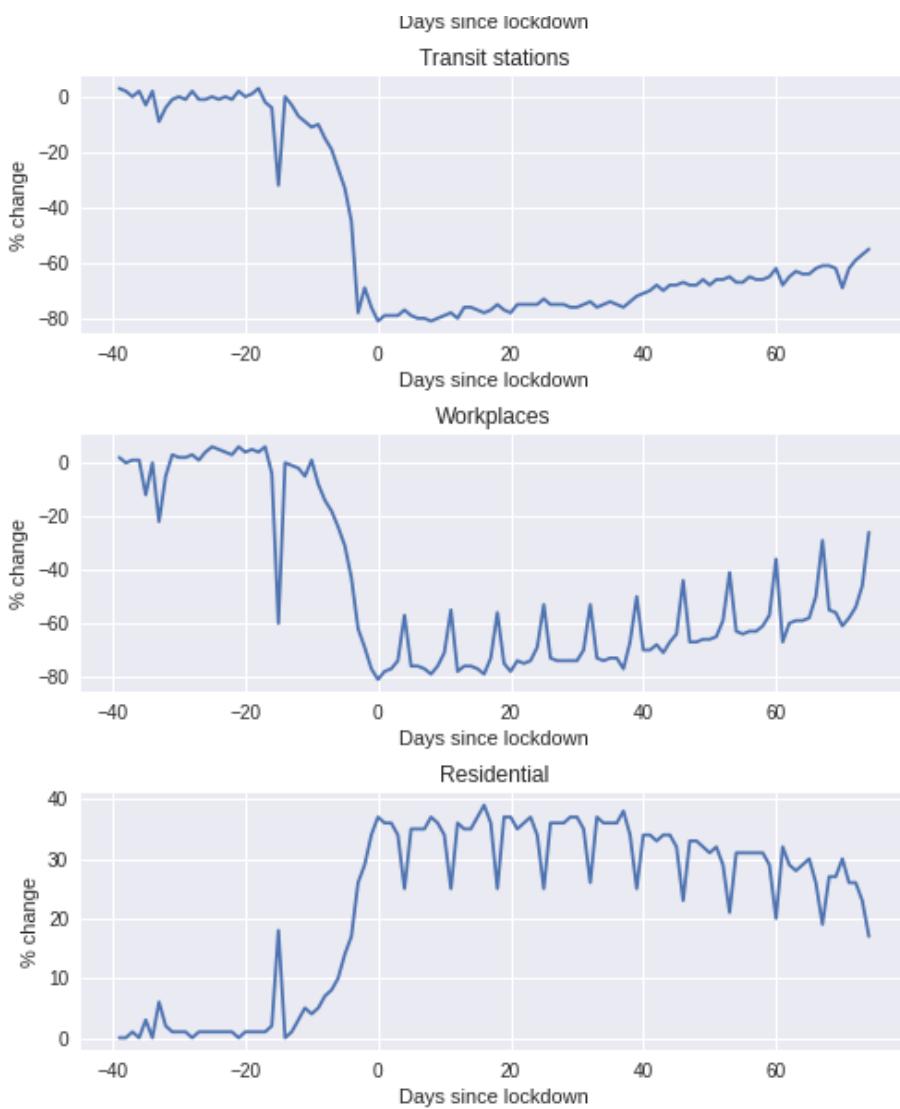
# Set the fifth subplot as active, and make the fifth plot.
plt.subplot(6, 1, 5)
plt.plot(x, y_workplaces_percent_change_from_baseline)
plt.title('Workplaces')
plt.xlabel('Days since lockdown')
plt.ylabel('% change')

# Set the final subplot as active, and make the final plot.
plt.subplot(6, 1, 6)
plt.plot(x, y_residential_percent_change_from_baseline)
plt.title('Residential')
plt.xlabel('Days since lockdown')
plt.ylabel('% change')

# Show the figure.
plt.show()

```





TODO Q10: Answer the following question: Do the results seem to satisfy “common sense”? Explain, citing specific data from your plot to support your answer.

TODO Q11: In the [Calibrate Region](#) checklist, Google suggests a number of reasons why the data might *not* be useful for understanding the effect of COVID-19-related lockdowns, or why the data might be misleading. For the location you have chosen, briefly answer all of the questions in that checklist. Based on your answers, do you think there are any serious problems associated with using this data for understanding user mobility changes due to COVID-19?

In [0]:

```
#Added code block so that answers can be easily distinguished
```

Q10 Response For the most part data seems to make sense. We can notice that Grocery and Pharmacy purchases are reaching normal stages. We can also see 10 peaks for Residential and workspaces signifying the weekends. So, we see a decline in people staying in residential areas on weekends. But these peaks don't make any sense for Workspaces - maybe they have been asked to report once a week.

Question 11 Response Let's try to break down our data with the given 5 checklist.

Note: The baseline is calibrated between Jan 3rd and Feb 6th of 2020, so all the percentage changes are with respective data.

Did anything significant happen between Jan 3 and Feb 6, 2020?

The main significant event that happened is the annual Indian festival called "Pongal" which lasts for one week, in which people go to their hometowns and celebrate it.

How do park visitors change from January to now?

Climate is more comfortable in January and February compared to the summer months, which accounts for

more frequent park visits during those months. We also generally don't find park data available to compare this trend with the previous years.

How much more time do you think people will spend in residential places?

From the first lockdown, we know that schools are closed but some offices have used different techniques to limit capacity which are not openly known. At the same time few medical staff had to work double shifts which will impact the time spent in residential areas, which we need to take into account.

How might types of work affect the mobility changes on weekdays or weekends?

Due to the travel restrictions, online food delivery systems like swiggy started delivering food with drones in certain limited areas which is impossible to know the reason for reduced congestion in those areas by looking into the dataset.

How well is your region represented?

The data didn't consider a lot of mini grocery stores that people still go to as they are not recognised by Google maps, which account for a considerable chunk in grocery purchases. It also didn't account for places like dockyard and air bases that people are going where phones are not allowed. For instance, there is a museum in front of the dockyard where most phones are turned off and back on before leaving and entering the base, so this might be considered a recreational visit and negatively affect the visitors to workplaces.

All in all, the data has a lot of problems which requires a lot of domain knowledge to filter out the data. Even, this will not be sufficient as there are flaws in baseline data as well. Thus, a model trained on this data will not be able to predict accurate results.

In [0]: